

Flume

Apache Flume, what is it?

- ▶ It is a distributed data collection service that gets flows of data (like logs) from their source and aggregates them to where they have to be processed.
- ▶ Goals: reliability, scalability, extensibility, manageability.

Challenges

- ▶ Physically Distributed Data Sources

Across servers, Data Centers, Geographies, Organization Boundaries, Technologies...

- ▶ Continuous Data Production

Every two days now we create as much information as we did from the dawn of civilization up until 2003"

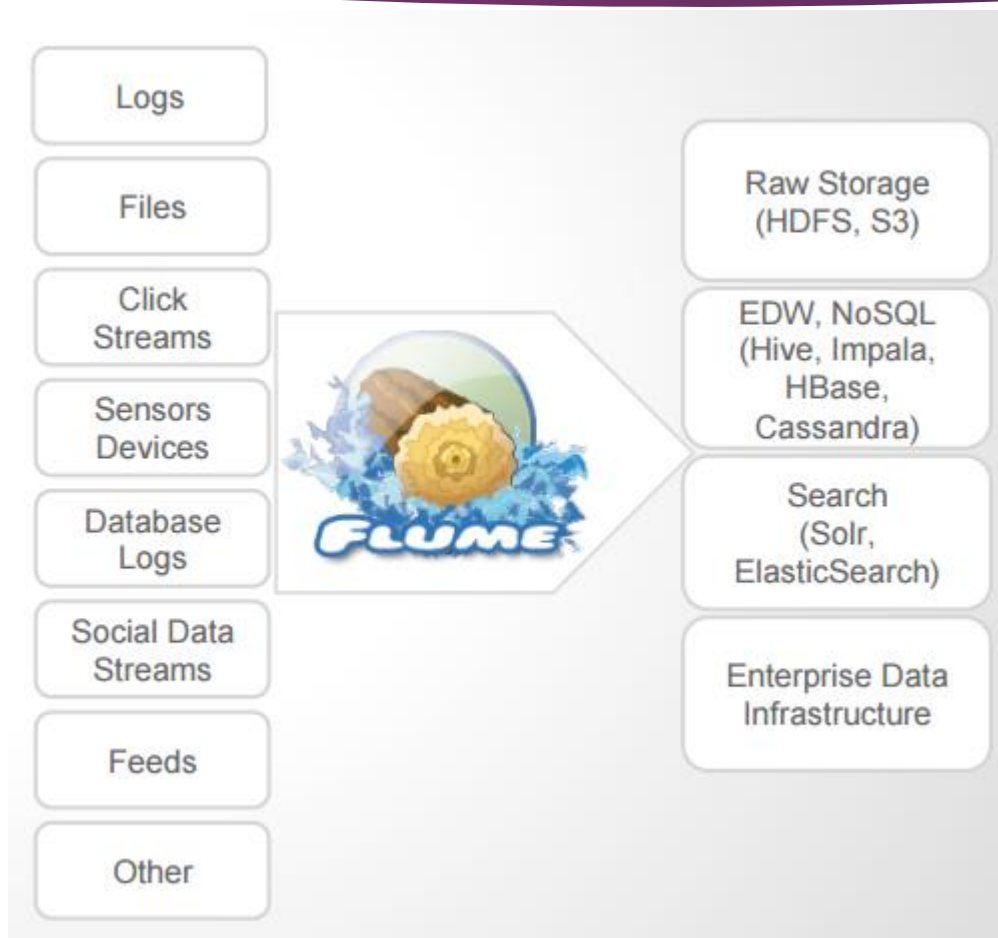
- Eric Schmidt, 2010

• Weather • Traffic • Automobiles • Trains • Airplanes • Geological/Seismic • Oceanographic • Smart Phones • Health Accessories • Medical Devices • Home Automation • Digital Cameras • Social Media • Geolocation • Shop Floor Sensors • Network Activity • Industry Appliances • Security/Surveillance • Server Workloads • Digital Telephony • Bio-simulations...

- ▶ Continuously evolving data structures and semantics...

CSV, TSV, JSON, XML, Binary, Custom Formats...

Intro



History

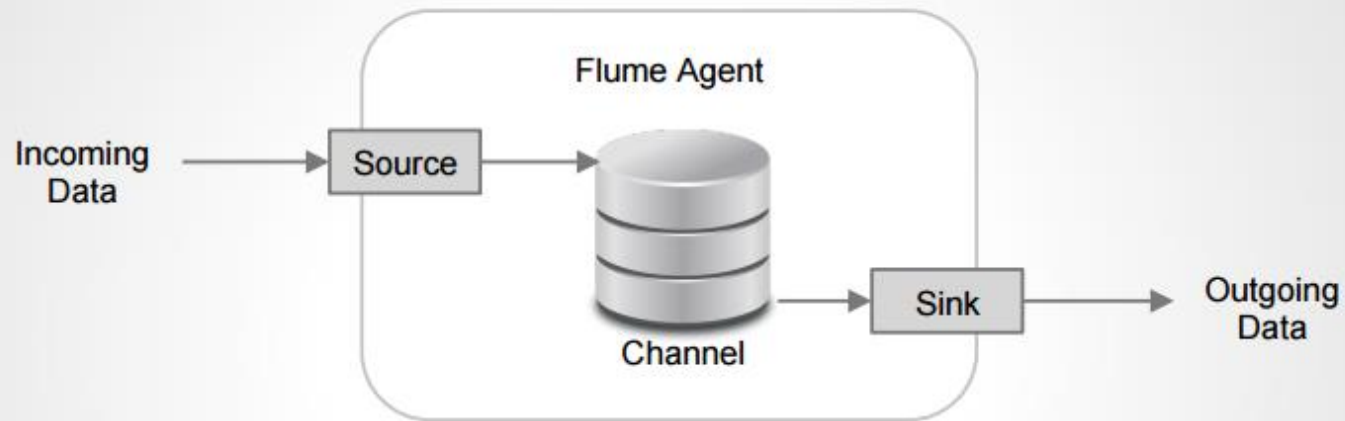
- ▶ Originally designed to be a log aggregation system by Cloudera Engineers
- ▶ Evolved to handle any type of streaming event data
- ▶ Low-cost of installation, operation and maintenance
- ▶ Highly customizable and extendable

Architecture

- ▶ Distributed Pipeline Architecture
- ▶ Optimized for commonly used data sources and destinations
- ▶ Built in support for contextual routing
- ▶ Fully customizable and extendable



Architecture



Source

- Accepts incoming Data
- Scales as required
- Writes data to Channel

Channel

- Stores data in the order received

Sink

- Removes data from Channel
- Sends data to downstream Agent or Destination

Core Concepts: Event

An Event is the fundamental unit of data transported by Flume from its point of origination to its final destination. Event is a byte array payload accompanied by optional headers.

- ▶ Payload is opaque to Flume
- ▶ Headers are specified as an unordered collection of string key-value pairs, with keys being unique across the collection
- ▶ Headers can be used for contextual routing

Core Concepts: Client

An entity that generates events and sends them to one or more Agents.


- ▶ Example
 - ▶ Flume log4j Appender
 - ▶ Custom Client using Client SDK (`org.apache.flume.api`)
- ▶ Decouples Flume from the system where event data is consumed from
- ▶ Not needed in all cases

Core Concepts: Agent

A container for hosting Sources, Channels, Sinks and other components that enable the transportation of events from one place to another.

- ▶ Fundamental part of a Flume *flow*
- ▶ Provides Configuration, Life-Cycle Management, and Monitoring Support for hosted components

Typical Aggregation Flow



$[Client]^+ \rightarrow Agent \rightarrow Agent^* \rightarrow Destination$

Core Concepts: Source

An active component that receives events from a specialized location or mechanism and places it on one or Channels.

- ▶ Different Source types:
 - ▶ Specialized sources for integrating with well-known systems. Example: Syslog, Netcat
 - ▶ Auto-Generating Sources: Exec, SEQ
 - ▶ IPC sources for Agent-to-Agent communication: Avro
- ▶ Require at least one channel to function

Core Concepts: Channel

A passive component that buffers the incoming events until they are drained by Sinks.

- ▶ Different Channels offer different levels of persistence:
 - ▶ Memory Channel: volatile
 - ▶ File Channel: backed by WAL implementation
 - ▶ JDBC Channel: backed by embedded Database
- ▶ Channels are fully transactional
- ▶ Provide weak ordering guarantees
- ▶ Can work with any number of Sources and Sinks.

Core Concepts: Sink

An active component that removes events from a Channel and transmits them to their next hop destination.

- ▶ Different types of Sinks:
 - ▶ Terminal sinks that deposit events to their final destination. For example: HDFS, HBase
 - ▶ Auto-Consuming sinks. For example: Null Sink
 - ▶ IPC sink for Agent-to-Agent communication: Avro
- ▶ Require exactly one channel to function

Configuration

agent1.properties:

```
# Active components
agent1.sources = src1
agent1.channels = ch1
agent1.sinks = sink1

# Define and configure src1
agent1.sources.src1.type = netcat
agent1.sources.src1.channels = ch1
```

```
agent1.sources.src1.bind = 127.0.0.1
agent1.sources.src1.port = 10112
```

```
# Define and configure sink1
agent1.sinks.sink1.type = logger
agent1.sinks.sink1.channel = ch1
```

```
# Define and configure ch1
agent1.channels.ch1.type = memory
```

Configuration

- ▶ A configuration file can contain configuration information for many Agents
- ▶ Only the portion of configuration associated with the name of the Agent will be loaded
- ▶ Components defined in the configuration but not in the active list will be ignored
- ▶ Components that are misconfigured will be ignored
- ▶ Agent automatically reloads configuration if it changes on disk

An Example

- ▶ Download flume
- ▶ Set path etc...
- ▶ Create a config file (say flume1.conf in ~/flumetests/ dir) with the below content –

```
myagent.sources = filesrc
myagent.sinks = logsink
myagent.channels = c1

# Describe/configure the source
myagent.sources.filesrc.type = spooldir
myagent.sources.filesrc.spoolDir = /home/hdtester

# Describe the sink
myagent.sinks.logsink.type = logger

# Use a channel which buffers events in file
myagent.channels.c1.type = file

# Bind the source and sink to the channel
myagent.sources.filesrc.channels = c1
myagent.sinks.logsink.channel = c1
```
- ▶ Run the agent

```
flume-ng agent -f flume1.conf -n myagent -Dflume.root.logger=INFO,console
```

One more example with HDFS + Console Sinks (Fan out)

File Name flume2.conf

```
myagent.sources = dirsrc
myagent.sinks = logsink hdfssink
myagent.channels = c1 c2

# Describe/configure the source
myagent.sources.dirsrc.type = spooldir
myagent.sources.dirsrc.spoolDir = /home/hdtester

# Describe the sink
myagent.sinks.logsink.type = logger

myagent.sinks.hdfssink.type = hdfs
myagent.sinks.hdfssink.hdfs.path = /flumecreated
myagent.sinks.hdfssink.hdfs.fileType = DataStream

# Use a channel which buffers events in file
myagent.channels.c1.type = file
myagent.channels.c1.checkpointDir = ../flume/file-
channel1/checkpoint
myagent.channels.c1.dataDirs = ../flume/file-
channel1/data
```

```
myagent.channels.c2.type = file
myagent.channels.c2.checkpointDir = ../flume/file-
channel2/checkpoint ##without this there would be a
lock issue with channel1
myagent.channels.c2.dataDirs = ../flume/file-
channel2/data ##without this there would be a lock
issue with channel1
```

```
# Bind the source and sink to the channel
myagent.sources.dirsrc.channels = c1 c2
myagent.sinks.logsink.channel = c1
myagent.sinks.hdfssink.channel = c2
```

```
flume-ng agent -f flume2.conf -n myagent -Dflume.root.logger=INFO,console
```

An example with 2 sources, 1 sink

File Name flume3.conf

```
myagent.sources = source1 source2
myagent.sinks = sink1
myagent.channels = channel1

myagent.sources.source1.type = exec
myagent.sources.source1.command = tail -f FILE1
myagent.sources.source1.channels = channel1

myagent.sources.source2.type = exec
myagent.sources.source2.command = tail -f FILE2
myagent.sources.source2.channels = channel1

myagent.channels.channel1.type = memory

myagent.sinks.sink1.type = file_roll
myagent.sinks.sink1.batchSize = 2
myagent.sinks.sink1.channel = channel1
myagent.sinks.sink1.sink.directory = /home/hdtester/flumetests
myagent.sinks.sink1.sink.rollInterval = 30
myagent.sinks.sink1.sink.serializer = TEXT
```

In another window, simulate writing into two files, named FILE1 and FILE2

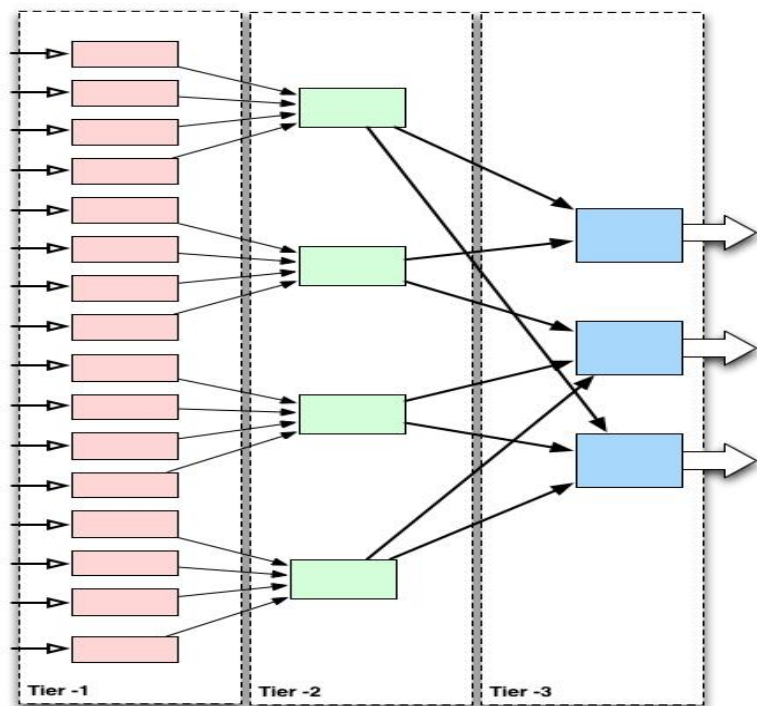
```
while true; do    echo "File 1 `date`" >> FILE1;
echo "File 2 `date`" >> FILE2;  sleep `shuf -i 1-5 -n
1`;done
```

```
flume-ng agent -f flume3.conf -n myagent -Dflume.root.logger=INFO,console
```

Add a HDFS sink to previous example

```
...  
myagent.sinks = sink1 sink2  
myagent.channels = channel1 channel2  
...  
myagent.channels.channel2.type = memory  
...  
myagent.sinks.sink2.type = hdfs  
myagent.sinks.sink2.hdfs.path = /flumecreated/h1  
myagent.sinks.sink2.hdfs.fileType = DataStream  
myagent.sinks.sink2.channel = channel2
```

Configuration



Typical Deployment

- ▶ All agents in a specific tier could be given the same name
- ▶ One configuration file with entries for three agents can be used throughout