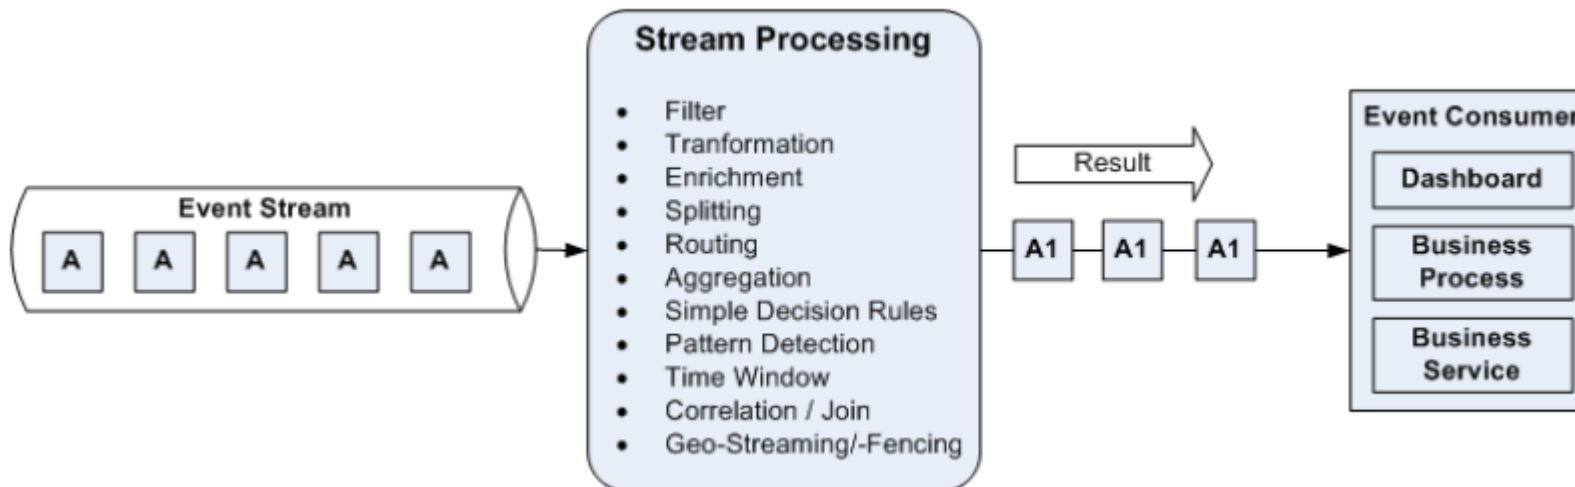


Storm and Trident

ixAT Solutions – ixatsolutions@gmail.com

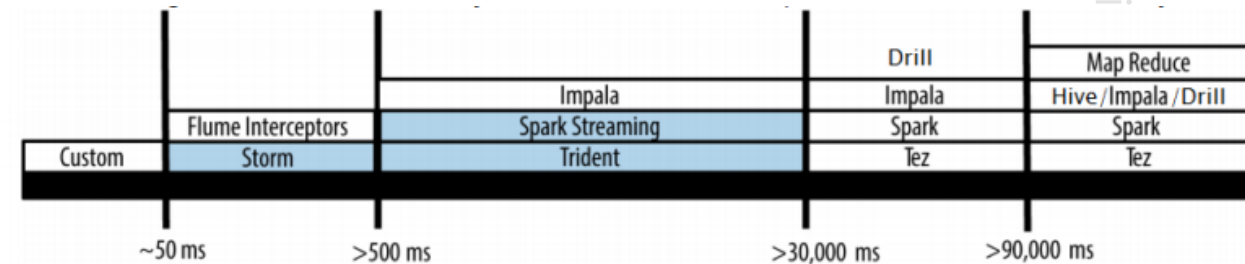
Stream Processing

- Infrastructure for continuous data processing.
- Computational model can be as general as MapReduce but with the ability to produce low-latency results
- Data collected continuously is naturally processed continuously aka. Event Processing / Complex Event Processing (CEP)



Processing Models

- ▶ Batch Processing
 - ▶ Familiar concept of processing data en masse
 - ▶ Generally incurs a high-latency
- ▶ (Event-) Stream Processing
 - ▶ A one-at-a-time processing model
 - ▶ A datum is processed as it arrives
 - ▶ Sub-second latency
 - ▶ Difficult to process state data efficiently
- ▶ Micro-Batching
 - ▶ A special case of batch processing with very small batch sizes (tiny)
 - ▶ A nice mix between batching and streaming
 - ▶ At cost of latency
 - ▶ Gives stateful computation, making windowing an easy task




Choices

- ▶ Latency
 - ▶ Is performance of streaming application paramount
- ▶ Development Cost
 - ▶ Is it desired to have similar code bases for batch and stream processing
=>lambda architecture
- ▶ Message Delivery Guarantees
 - ▶ Is there high importance on processing every single record, or is some normal amount of data loss acceptable
- ▶ Process Fault Tolerance
 - ▶ Is high-availability of primary concern

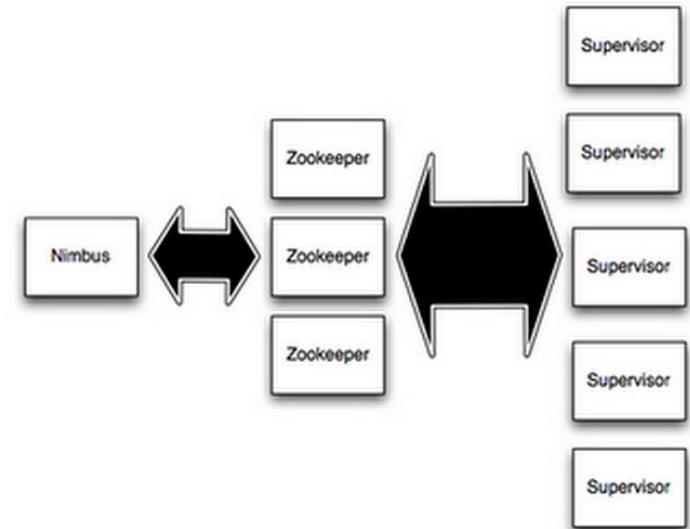


Storm Overview

- ▶ DAG Processing of never ending streams of data
 - ▶ Open Sourced: <https://github.com/nathanmarz/storm/wiki>
 - ▶ Used at Twitter plus > 24 other companies
 - ▶ Reliable - At Least Once semantics
 - ▶ Think MapReduce for data streams
 - ▶ Java / Clojure based
 - ▶ Bolts in Java and 'Shell Bolts'
 - ▶ Not a queue, but usually reads from a queue.
- ▶ Related:
 - ▶ S4, CEP, InfoStreams, HStreaming,...
- ▶ Compromises
 - ▶ Static topologies & cluster sizing – Avoid messy dynamic rebalancing
 - ▶ Nimbus – SPOF
- ▶ Strong Community Support, commercial support from InfoChimps

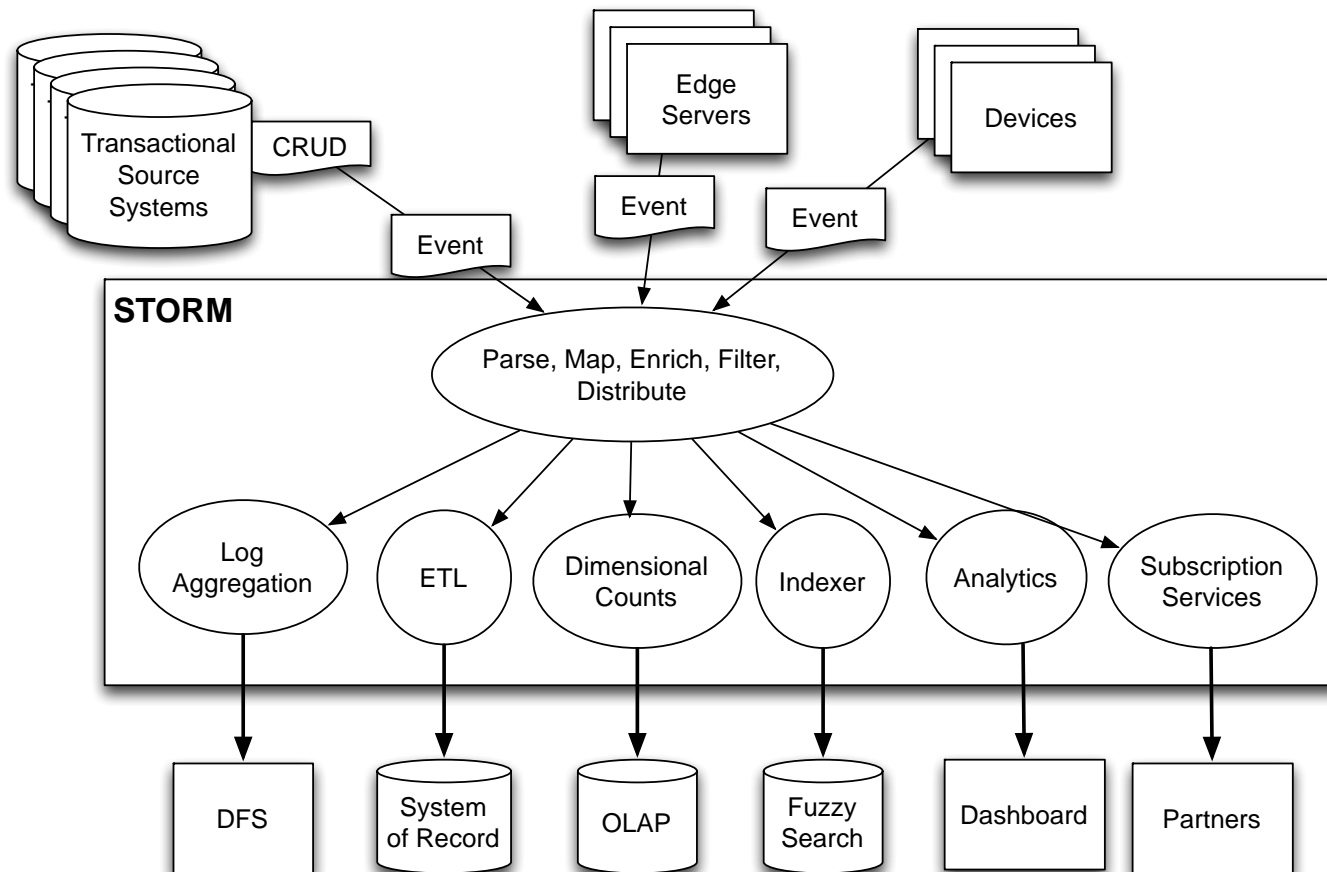


Storm Concepts

- ▶ Cluster
 - ▶ Supervisor
 - ▶ Worker
- ▶ Topology
- ▶ Streams
- ▶ Spout
- ▶ Bolt
- ▶ Tuple
- ▶ Stream



Storm & Big Data Patterns



What is Trident exactly?

- ▶ Trident is
 - ▶ Abstraction on top of Storm Infrastructure
 - ▶ Topology
 - ▶ Divided into streams
 - ▶ Optimized
 - ▶ DRPC compatible
 - ▶ Added to Storm 0.8.0 – Aug 2012
- ▶ Simplifies building topologies
- ▶ Core data model is the stream
 - ▶ Processed as a series of batches (micro-batches)
 - ▶ Stream is partitioned among nodes in cluster
- ▶ 5 kinds of operations in Trident
 - ▶ Operations that apply locally to each partition and cause no network transfer
 - ▶ Repartitioning operations that don't change the contents
 - ▶ Aggregation operations that do network transfer
 - ▶ Operations on grouped streams
 - ▶ Merges and Joins

Comparision

	Core Storm	Storm Trident	Spark Streaming
Community	> 100 contributors	> 100 contributors	> 280 contributors
Adoption	***	*	*
Language Options	Java, Clojure, Scala, Python, Ruby, ...	Java, Clojure, Scala	Java, Scala Python (coming)
Processing Models	Event-Streaming	Micro-Batching	Micro-Batching Batch (Spark Core)
Processing DSL	No	Yes	Yes
Stateful Ops	No	Yes	Yes
Distributed RPC	Yes	Yes	No
Delivery Guarantees	At most once / At least once	Exactly Once	Exactly Once
Latency	sub-second	seconds	seconds
Platform	Storm Cluster, YARN	Storm Cluster, YARN	YARN, Mesos Standalone, DataStax EE